



A detailed botanical illustration featuring several different plant species. On the left, a tall, slender plant with a dark, textured, rounded seed head at the top. Next to it is a plant with a green stem and several pointed, serrated leaves. In the center, a plant with a green stem and a small, elongated, textured seed head. To the right of that, a plant with a green stem and a large, yellow, elongated seed head. Further right, a plant with a green stem and a small, elongated, textured seed head. On the far right, a plant with a green stem and several pointed, serrated leaves. The background is a light cream color.

Apprentissage dans les Systèmes Naturels et Artificiels

Chaîne de traitement statistique / ML Workflow 2



Sélection d'attributs/
Feature Selection

Feature selection Vs Dimensionality reduction

La sélection d'attributs vise à réduire le nombre d'attributs dans le jeu d'entraînement.

Pourquoi ?

- Pour faciliter le travail de l'algorithme
- Pour trier les attributs importants dans l'analyse supervisée et obtenir un modèle plus simple (rasoir d'occam)

La réduction de dimension modifie les attributs (combinaison).

La sélection d'attributs en élimine certains sans les changer.

Feature Selection

Techniques simples (simplistes?)

- Supprimer les attributs avec la plus faible variance
- Utilisation d'un test statistique (χ^2) entre un attribut et les labels et ne conserver que les n meilleurs attributs ou les n%.

Méthodes qui ont été (trop) populaires:

- RFE (Recursive Feature Elimination): on itère le modèle sur un jeu chaque fois réduit de l'attribut au poids le moindre.

Feature Selection

Méthodes meilleures?

- Utiliser un modèle sélectif (e.g. L1 regression) avant d'entraîner le modèle supervisée
- Utiliser un modèle supervisée sélectif nativement



Le plus important?

- À chaque jeu d'entraînement, son ensemble d'attributs sélectionnés
- Il faut analyser la stabilité de la sélection (Nogueira 2018)
- Au final on peut discuter les attributs les plus fréquemment sélectionnés...

Feature selection

Autres pratiques:

- Rejeter les attributs comportant des données manquantes (plutôt que les échantillons)
- Rejeter un attribut s'il est fortement corrélé à un autre



Extraction d'attributs/ Feature Extraction

Feature extraction

Feature extraction $\neq$ Feature selection !!

Le terme « feature extraction » est employé pour la création d'attributs employables par les algorithmes ML à partir de textes ou d'images

Textes : exemple du Bag of Words (tokenizing), un échantillon est un document, un attribut la fréquence d'un « token » et la matrice celle d'un corpus de documents.



Feature extraction

Images :

- VP-Tree : résumé dans un 32-bit integer
- Proportion des valeurs RGB
- Kaze, etc.

Objectif principal : trouver des images/objets similaires

Inconvénient : réducteur et souvent inutile pour l'analyse supervisée



Réduction de
dimension

ACP/PCA

Analyse en composantes principales

Aide à formuler des hypothèses... n'est pas une fin en soi


Principe : nouveau repère de combinaisons linéaires des attributs dont on ne retiendra que les n premiers axes appelés les composantes principales qui résument un maximum de variance.

Intérêt : visualisation, aide à la compréhension des interactions, isoler un bruit, **feature selection**


Inconvénients : difficile à interpréter, perte d'information

La PCA est non supervisée

ACP/PCA



ACP/PCA







Feature agglomeration

Basé sur le clustering hiérarchique

Groupe ensemble les attributs qui se comportent de manière similaire.

Feature agglomeration





Sur-ajustement(s)

Terminologie

- " Overfitting " is traditionally defined as training some flexible representation so that it memorizes the data but fails to predict well in the future.
- Autre définition : over-representing the performance of systems

Sur-ajustement classique

Traditional overfitting: Train a complex predictor on too-few examples.

Remedy:

- Hold out pristine/very good examples for testing.
- Use a simpler predictor.
- Get more training examples.
- Integrate over many predictors.
- Reject this.

Sur-ajustements subtiles

Parameter tweak overfitting: Use a learning algorithm with many parameters. Choose the parameters based on the test set performance.

For example, choosing the features so as to optimize test set performance can achieve this.

Sur-ajustements subtiles

Brittle measure: Use a measure of performance which is especially brittle to overfitting.

Examples: "entropy", "mutual information", and leave-one-out cross validation are all surprisingly brittle.

Sur-ajustements subtiles

Bad statistics: Misuse statistics to overstate confidences.

One common example is pretending that cross validation performance is drawn from an i.i.d. gaussian, then using standard confidence intervals. Cross validation errors are not independent. Another standard method is to make known-false assumptions about some system and then derive excessive confidence.

Sur-ajustements subtiles

Choice of measure: Choose the best of Accuracy, error rate, (A)ROC, F1, percent improvement on the previous best, percent improvement of error rate, etc.. for your method. For bonus points, use ambiguous graphs.

This is fairly common and tempting.

Remedy: Use canonical performance measures. For example, the performance measure directly motivated by the problem.

Sur-ajustements subtiles

Human-loop overfitting: Use a human as part of a learning algorithm and don't take into account overfitting by the entire human/computer interaction.

This is subtle and comes in many forms. One example is a human using a clustering algorithm (on training and test examples) to guide learning algorithm choice.

Remedy: Make sure test examples are not available to the human.

Sur-ajustements subtiles

Data set selection: Chose to report results on some subset of datasets where your algorithm performs well.

The reason why we test on natural datasets is because we believe there is some structure captured by the past problems that helps on future problems. Data set selection subverts this and is very difficult to detect.

Remedy: Use comparisons on standard datasets. Select datasets without using the test set. Good Contest performance can't be faked this way.

Sur-ajustements subtiles

Reproblemenging: Alter the problem so that your performance improves.

Remedy: For example, take a time series dataset and use cross validation. Or, ignore asymmetric false positive/false negative costs. This can be completely unintentional, for example when someone uses an ill-specified UCI dataset.

Remedy: Discount papers which do this. Make sure problem specifications are clear.

Sur-ajustements subtiles

Old datasets: Create an algorithm for the purpose of improving performance on old datasets.

After a dataset has been released, algorithms can be made to perform well on the dataset using a process of feedback design, indicating better performance than we might expect in the future. Some conferences have canonical datasets that have been used for a decade...



Remedy: Prefer simplicity in algorithm design. Weight newer datasets higher in consideration. Making test examples not publicly available for datasets slows the feedback design process but does not eliminate it.

Sur-ajustements subtiles

Overfitting by review: 10 people submit a paper to a conference. The one with the best result is accepted. This is a systemic problem which is very difficult to detect or eliminate. We want to prefer presentation of good results, but doing so can result in overfitting.


Remedy:

- Be more pessimistic of confidence statements by papers at high rejection rate conferences.



Conclusion

Conclusion



Conclusion

- Préciser son attente
- Concevoir sa chaîne de traitement
- Ne pas sur-interpréter les résultats

- le jeu d'entraînement est un bac à sable où tout est permis
- Le jeu de test doit être rigoureusement tenu à l'écart

- Trouver un équilibre pour éviter le sur-ajustement et obtenir une meilleure généralisation